

Screaming Channels on Bluetooth Low Energy – ACSAC Artifact Evaluation

Pierre AYOUB

August 28, 2024

Contents

1	Abstract	1
2	Overview	2
3	Setup	3
3.1	Download	3
3.2	Installation	3
4	Reproducing the attacks	4
4.1	Attacking A_7	4
4.2	Attacking A_9	7
5	Clean	9
6	Conclusion	9

1 Abstract

This artifact consists of software and data to reproduce the main results of section 8.5 of our paper. The software is composed of Python source code and a C library. Everything is automated inside a Docker container, and dependencies and manual installation are referenced under the main project web page. The data mainly consists of Numpy arrays containing complex numbers representing I/Q data.

2 Overview

This guide will help you reproduce the main results of our paper. Note that this is a small part of our project. The full project can be found on GitHub at [pierreay/screaming_channel_ble](https://github.com/pierreay/screaming_channel_ble). A web version of this guide can also be found at screaming_channels_ble/docs/demo_20240828_acsac/README.org. We have published in open-access all of our data in the repository located at pierreay/screaming_channels_data.

To fully reproduce this attack, one would have to first acquire a dataset by:

1. Flashing two firmware (the attacker dongle and the victim board) and installing the radio tooling.
2. Setup the experimental setup, including the attacker host computer, the radio, the antenna and amplifiers, the dongle, the victim dongle.
3. Performing a dataset collection (up to several days).
4. Post-processing the dataset to extract the small AES leak amplitude trace from the large RAW I/Q signal.

As you imagine, this is fairly complex and long without prior experience. Hence, in this guide, we provides two representative pre-processed datasets on which we can complete the second part of the attack. The A_x refer to the paper notation, while the long name refer to our data repository notation:

A_7 240207_1-leak-pairing-10cm-anechoic-2.533e9-8e6_raw

A_9 240429_highdist_2lna_highgain_norep

Hence, we uploaded two representative and small datasets on Zenodo:

Ayoub (2024) Screaming Channels on Bluetooth Low Energy, Zenodo. DOI: 10.5281/zenodo.13384278. URL: <https://zenodo.org/records/13384278>

For both datasets, we will execute the following steps:

1. Create a profile (*i.e.*, a template) from a training subset to learn the leakage model.
2. Leveraging the profile, attack using a Template Attack on an attack subset.

While it is not mandatory, the docs/demo_20231014/README.org demo give a look at visualizing how traces and the AES leakage looks like.

3 Setup

First, we will setup necessary tools, downloading the code and the data.

3.1 Download

Clone the "Screaming Channels on Bluetooth Low Energy" repository (~ 10 MB) in the directory of your choice:

```
git clone https://github.com/pierreay/screaming_channels_ble/
```

Moreover, manually download the two aforementioned datasets located at zenodo.org/records/13384278 which have been publicly uploaded on Zenodo (less than 2 GB).

3.2 Installation

We will setup a temporary Docker container for reproducing the attacks. Ensure that both Docker and its buildx builder are installed.

First, move the downloaded datasets inside the Docker folder:

```
cd screaming_channels_ble/docs/demo_20240828_acsac
mv -t . /PATH/T0/240207_1-leak-pairing-10cm-anechoic-2.533e9-8e6_raw.tar.bz2
mv -t . /PATH/T0/240429_highdist_2lna_highgain_norep.tar.bz2
```

Second, download and initialize the Docker image leveraging the `Dockerfile`:

```
make build
```

It will download around 2 GB and decompress the datasets that have been downloaded previously inside the container. The image and containers will be cleaned up at the end, no files will be left or modified on the host system.

If you get an error like the following one:

```
cp: cannot stat '/host/*.tar.bz2': No such file or directory
```

It means that the two datasets archives were not moved correctly to the Dockerfile directory and that the Docker container cannot access them.

4 Reproducing the attacks

The Docker is now ready to be use. In the following, we assume that the reader is connected to the Docker container through SSH:

```
make shell
```

For the following command, the `$SC_SRC` variable is set to the path of the `screaming_channels_ble/src` directory, while the `$DATASET` variable will be set to the path of the currently analyzed dataset.

4.1 Attacking A_7

This dataset correspond to the scenario with the non-instrumented firmware in the anechoic box.

Set the path to the dataset:

```
export DATASET="/240207_1-leak-pairing-10cm-anechoic-2.533e9-8e6_raw"
```

First, create the profile:

```
"${SC_SRC}/attack.py" --dataset-path "${DATASET}" --custom-dtype --plot \
    --save-images --norm --comptype AMPLITUDE \
    --num-traces 16384 --start-point 1000 --end-point 1500 \
    profile --pois-algo r --num-pois 1 --poi-spacing 1 \
    --variable p_xor_k --align
```

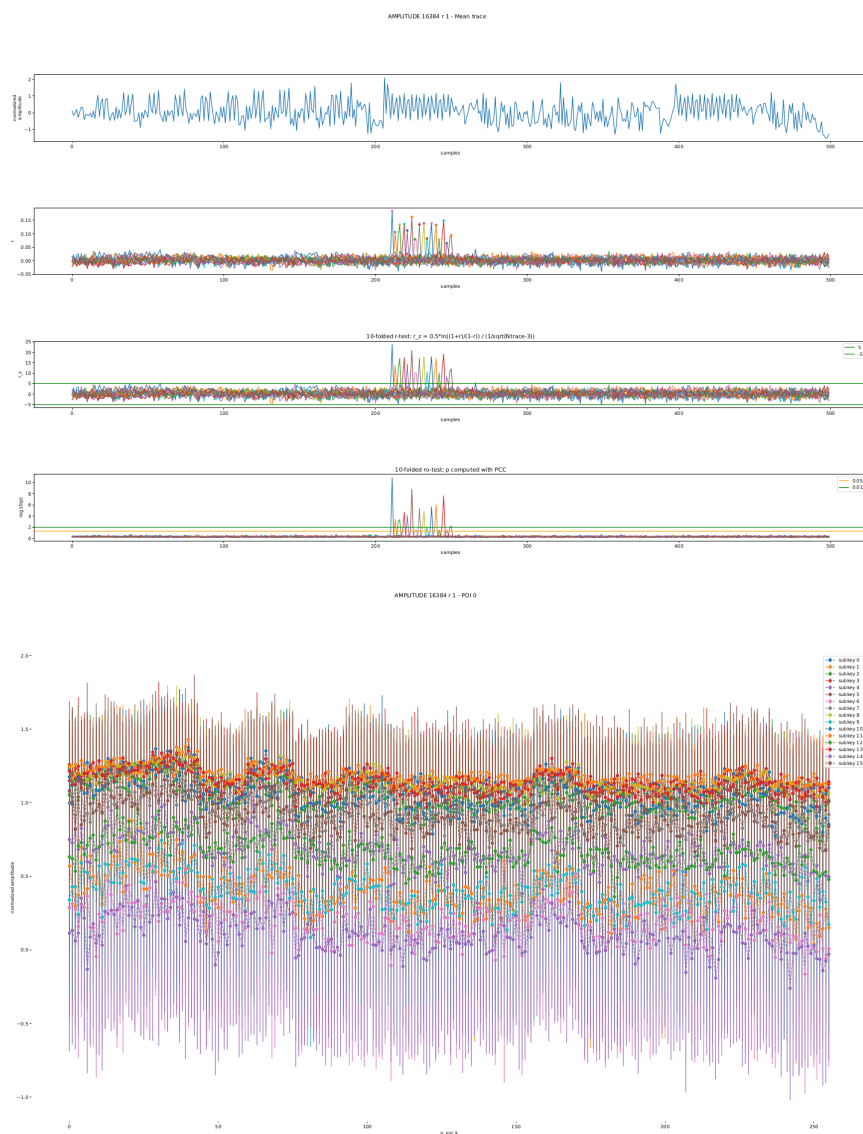
Where non-obvious options means:

- custom-dtype** Use a custom `dtype` instead of the default Numpy complex `dtype`, allowing to save disk space.
- norm** Normalize the traces after loading.
- start-point / –end-point** Truncate the traces between those points.
- pois-algo** Use the k – fold ρ – test to find informative point of interests (PoIs).
- num-pois** Use only 1 PoI per subbytes.
- variable $\mathbf{p_{xor_k}}$** Use the $l = p \oplus k$ leakage variable, with p the plaintext and k the key.

At that point, a plot should open (if X11 forwarding is correctly configured with Docker) and the profile will be stored under `#{DATASET}/profile`. If no plot appears, check that the profile has correctly been created by running:

```
evince #{DATASET}/profile/*.pdf
```

You should obtains plots similar to the following ones:



We will now attack using the created profile:

```
"${SC_SRC}"/attack.py --dataset-path "${DATASET}" --custom-dtype --log \
    --plot --norm --bruteforce --comptype AMPLITUDE \
    --start-point 1000 --end-point 1500 --num-traces 14900 \
    attack --attack-algo pcc --profile "${DATASET}/profile" --num-pois 1 \
    --poi-spacing 1 --variable p_xor_k --align-attack --align-profile \
    --no-align-profile-avg
```

The options are fairly similar to the previous command, except that we also specify to bruteforce the key at the end, and to re-perform a step of trace alignment. You should see this output:

```
Best Key Guess: 7c 19 67 e4 41 65 da 6d f5 18 32 f7 c5 f1 13 3e
Known Key:      7f 18 67 e6 42 65 da 6f f5 1b 32 f4 c6 f3 10 3e
PGE:            002 001 000 003 001 000 000 003 000 003 000 003 002 002 001 000
HD:             002 001 000 001 002 000 000 001 000 002 000 002 002 001 002 000
SUCCESS:        0 0 1 0 0 1 1 0 1 0 1 0 0 0 0 1
CORRECT BYTES: 6
PGE MEAN:      1
PGE MEDIAN:    1
PGE MAX:       3
HD SUM:        16
```

```
Starting key ranking using HEL
results rank estimation
nb_bins = 512
merge = 2
Starting preprocessing
Clearing memory
min: 2^33.79125918
actual rounded: 2^34.59872304
max: 2^35.23080812
time enum: 0.343729 seconds
```

Which means that this attack leads to an estimated key rank of 2^{34} , just as A_7 in the paper.

With the additional bruteforce, which took 1,5 hours running natively on my desktop computer but a little more inside the Docker container, you should obtain:

```
Starting key enumeration using HEL
Assuming that we know two plaintext/ciphertext pairs
[...]
Starting preprocessing
current rank : 2^2.584962501
current rank : 2^4.459431619
[...]
```

current rank : 2^{33.13069245}

current rank : 2^{34.11663176}

KEY FOUND!!!

7f 18 67 e6 42 65 da 6f f5 1b 32 f4 c6 f3 10 3e

Which means we break the entire key!

4.2 Attacking A_9

This dataset correspond to the scenario with the non-instrumented firmware at a distance of 1.2 meter in an office environment. Since the steps are similar to the previous attack, we will not detail the commands and options.


First, create the profile:


```
export DATASET="/240429_highdist_2lna_highgain_norep/raw"
```

```
"${SC_SRC}/attack.py" --dataset-path "${DATASET}" --custom-dtype --plot \
--save-images --norm --comptype AMPLITUDE \
--num-traces 12000 --start-point 1100 --end-point 1650 \
profile --pois-algo r --num-pois 1 --poi-spacing 1 \
--variable p_xor_k --align
```

```
evince ${DATASET}/profile/*.pdf
```

You should obtains plots similar to the following ones:





Then, attack using the created profile:

```
"${SC_SRC}"/attack.py --dataset-path "${DATASET}" --custom-dtype --log \
--plot --norm --no-bruteforce --comptype AMPLITUDE \
--start-point 1020 --end-point 1570 --num-traces 40000 \
attack --attack-algo pcc --profile "${DATASET}/profile" --num-pois 1 \
--poi-spacing 1 --variable p_xor_k --align-attack --align-profile \
--no-align-profile-avg
```

You should see this output:

```
Best Key Guess:  a3  e3  69  9c  a9  2a  3f  46  f4  5f  11  5c  78  80  a6  4f
Known Key:      ad  ef  6d  92  a8  29  32  40  f4  56  17  52  79  8d  a6  46
PGE:            006  001  014  005  002  005  008  010  000  007  011  003  006  004  000  010
HD:             003  002  001  003  001  002  003  002  000  002  002  003  001  003  000  002
SUCCESS:        0    0    0    0    0    0    0    0    1    0    0    0    0    0    1    0
CORRECT BYTES:  2
PGE MEAN:       5
PGE MEDIAN:     5
PGE MAX:        14
HD SUM:         30
```

```
Starting key ranking using HEL
results rank estimation
nb_bins = 512
merge = 2
Starting preprocessing
Clearing memory
min: 2^55.02283544
actual rounded: 2^60.13102936
max: 2^62.56824053
time enum: 0.196843 seconds
```


Which means that this attack leads to an estimated key rank of 2^{60} , just as A_9 in the paper.

5 Clean

Finally, once the Docker container is exited, one may run the following to clean the image:

```
make clean
```

6 Conclusion

In this demonstration, we reproduced the two most important results of the paper regarding attack performances. All other claims can also be reproduced by leveraging our full dataset repository, including scripts and data. Regarding the protocol manipulation part of the attack, a better solution would be to have an online platform allowing to replicate and extend such attacks. We are looking into this as future work.